

LoLWinProbabilityModel

May 11, 2025

```
[ ]: import pandas as pd
      #import lol matches for model
      data = pd.read_csv('ranked_14_min_stats_zero_center.csv')
      data.head(50)
```

```
[13]: from sklearn.preprocessing import OneHotEncoder
      #onehotencode champion names
      encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
      champion_onehot = encoder.fit_transform(data[['championName']])
      #get champion column out
      champion_onehot_columns = encoder.get_feature_names_out(['championName'])

      champion_df = pd.DataFrame(champion_onehot, columns = champion_onehot_columns)

      copied_final_df_encoded_champions = pd.concat([data.reset_index(drop = True),
      ↪champion_df], axis = 1)
      copied_final_df_encoded_champions.drop('championName', axis = 1, inplace = True)
```

```
[ ]: copied_final_df_encoded_champions.columns
```

```
[17]: #get rid of unneeded columns
      copied_final_df_encoded_champions_dropped_columns =
      ↪copied_final_df_encoded_champions.drop(columns = [

      ↪          'goldPerMinute',
      ↪          'minionsKilled',
      ↪          'jungleMinionsKilled',
      ↪          'csPerMinute',
      ↪          'participantId',
      ↪          'teamPosition'],axis = 'columns')
```

```
[23]: #rename 100/200 to Blue/Red
copied_final_df_encoded_champions_dropped_columns['teamId'] =
↳ copied_final_df_encoded_champions_dropped_columns['teamId'].replace({100:
↳ 'Blue',

↳
200: 'Red'})
```

```
[25]: #rename horde kills to grubs
copied_final_df_encoded_champions_dropped_columns =
    copied_final_df_encoded_champions_dropped_columns.rename(columns =
    {'teamHordeKills_14': 'teamGrubKills'})
```

```
[27]: #aggregate stats for team based modeling
champion_names = [col for col in
    ↪copied_final_df_encoded_champions_dropped_columns if col.
    ↪startswith('championName_')]

stat_agg = {
    'firstBloodKill': 'max',
    'totalGold' : 'sum',
    'xp': 'sum',
    'level': 'mean',
    'platesTaken_14' : 'sum',
    'towersKilled_14' : 'sum',
    'teamGrubKills': 'max',
    'teamDragonKills_14': 'max',
    'wards_placed': 'sum',
    'roleGoldDiff': 'sum',
    'roleXpDiff': 'sum',
    'roleCsDiff' : 'sum',
    'roleKillDiff' : 'sum',
    'roleDeathsDiff' : 'sum',
    'roleAssistsDiff' : 'sum',
    'roleVisionDiff' : 'sum',
    'win': 'max'}

for col in champion_names:
    stat_agg[col] = 'sum'

copied_final_df_encoded_champions_dropped_columns_grouped_team =
    ↪(copied_final_df_encoded_champions_dropped_columns
    .
    ↪groupby(['match_id', 'teamId'])
    .agg(stat_agg)
    .reset index())
```

```
[ ]: #pivot so we can get one row per match comparing both teams side-by-side as
      ↳ blue vs red
pivoted = copied_final_df_encoded_champions_dropped_columns_grouped_team.
      ↳ pivot(index='match_id', columns='teamId')

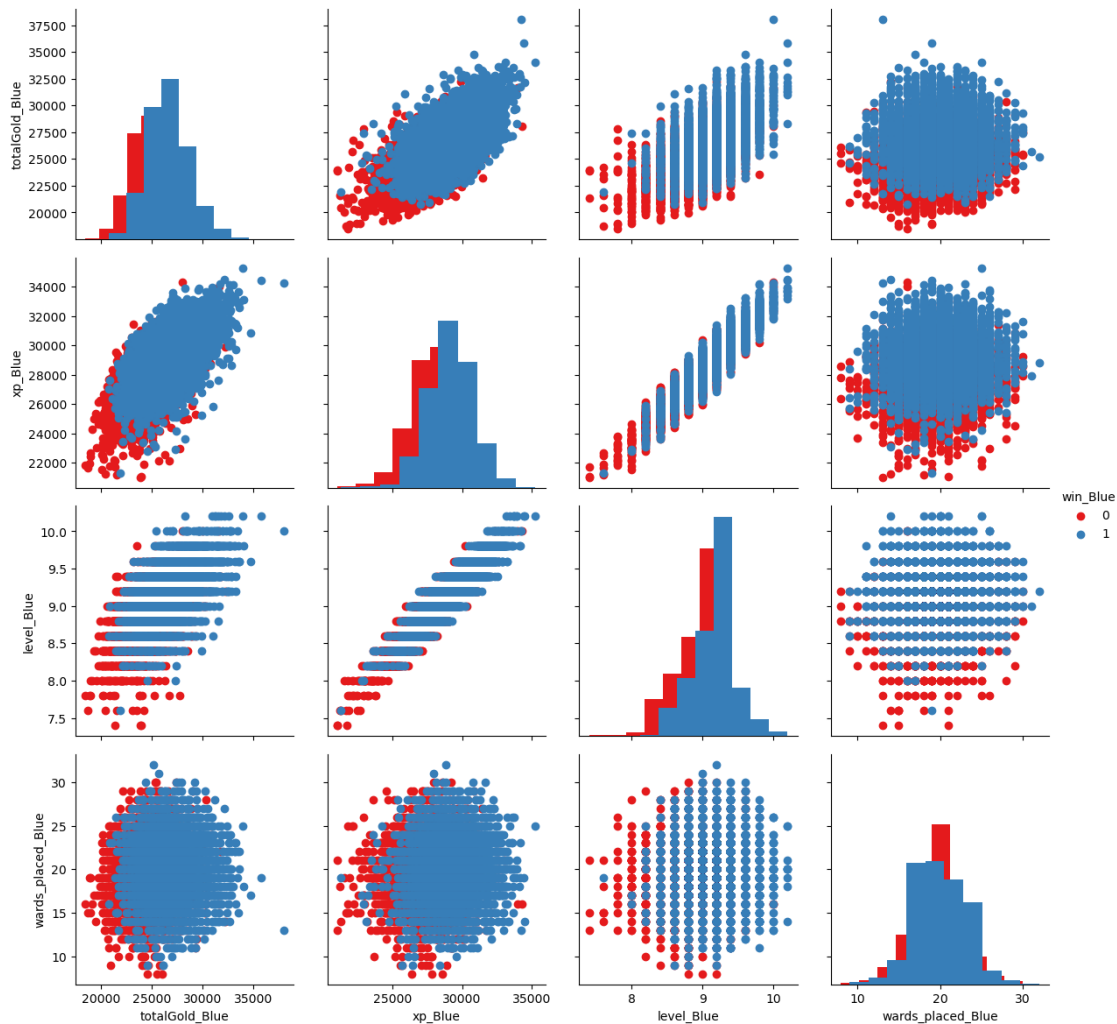
#flatten column names
pivoted.columns = ['{}_{}'.format(stat, team) for stat, team in pivoted.columns]
pivoted.reset_index(inplace=True)
pivoted
```

```
[33]: #rename role based headers since we are now looking at blue vs red team metrics
pivoted.columns = [
    col.replace("role", "") if col.startswith("role") else col
    for col in pivoted.columns]
```

```
[35]: #remove win_red since we are predicting blue win rate
pivoted = pivoted.drop(columns=['win_Red'], axis = 'columns')
```

```
[37]: #take a look at scatterplot matrix to visualize relationships between some of
      ↳ the features
import seaborn as sns
import matplotlib.pyplot as plt
pairgrid = sns.PairGrid(pivoted, vars =
    ↳ ['totalGold_Blue', 'xp_Blue', 'level_Blue', 'wards_placed_Blue'], hue =
    ↳ 'win_Blue', height = 3, palette = 'Set1')
pairgrid.map_diag(plt.hist)
pairgrid.map_offdiag(plt.scatter)
pairgrid.add_legend()
```

```
[37]: <seaborn.axisgrid.PairGrid at 0x156db2b10>
```



```
[39]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
import numpy as np

#split data into predictors and response
X = pivoted.drop(columns=['match_id', 'win_Blue'])
y = pivoted['win_Blue']

#initialize the model
model = RandomForestClassifier(n_estimators=100, random_state=42)

#5-fold cross-validation with ROC AUC scoring
scores = cross_val_score(model, X, y, cv=5, scoring='roc_auc')

#Output the results
```

```
print("ROC AUC scores from each fold:", scores)
print("Mean ROC AUC:", np.mean(scores).round(4))
```

ROC AUC scores from each fold: [0.84582887 0.85099096 0.84066587 0.85403146
0.8376285]

Mean ROC AUC: 0.8458

```
[41]: #fit the model
model.fit(X, y)

#get feature importances
importances = model.feature_importances_
feature_names = X.columns

#get indices of top 20 features, sorted descending
top_indices = importances.argsort()[-20:][::-1]

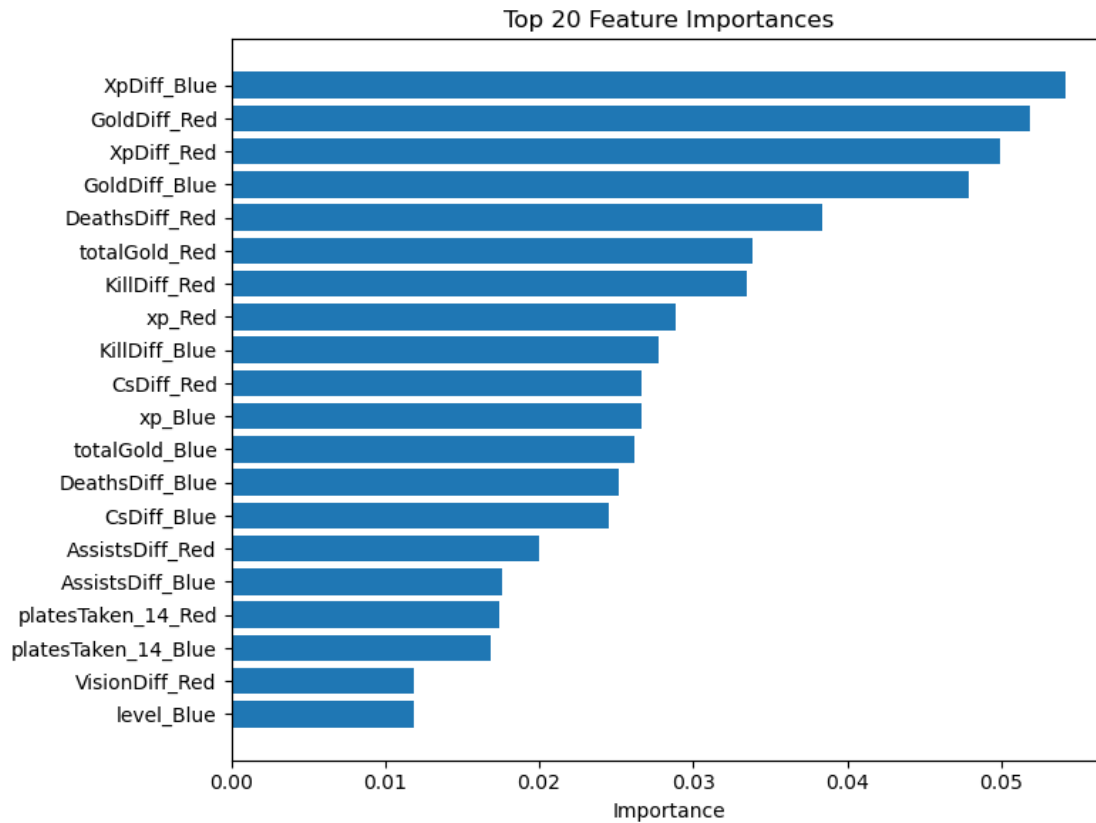
#print top 20 features from most to least important
for i in top_indices:
    print(f"{feature_names[i]}: {importances[i]:.4f}")
```

XpDiff_Blue: 0.0542
GoldDiff_Red: 0.0518
XpDiff_Red: 0.0500
GoldDiff_Blue: 0.0479
DeathsDiff_Red: 0.0384
totalGold_Red: 0.0339
KillDiff_Red: 0.0335
xp_Red: 0.0289
KillDiff_Blue: 0.0278
CsDiff_Red: 0.0266
xp_Blue: 0.0266
totalGold_Blue: 0.0262
DeathsDiff_Blue: 0.0252
CsDiff_Blue: 0.0245
AssistsDiff_Red: 0.0200
AssistsDiff_Blue: 0.0176
platesTaken_14_Red: 0.0174
platesTaken_14_Blue: 0.0168
VisionDiff_Red: 0.0119
level_Blue: 0.0118

```
[79]: #feature importance
import numpy as np

importances = model.feature_importances_
indices = np.argsort(importances)[-20:] # top 20
plt.figure(figsize=(8, 6))
```

```
plt.barh(np.array(X.columns)[indices], importances[indices])
plt.title("Top 20 Feature Importances")
plt.xlabel("Importance")
plt.tight_layout()
plt.show()
```



```
[43]: from sklearn.model_selection import GridSearchCV

#number of trees in forest, depth of each tree, and min samples required to
↳split the node
param_grid = {'n_estimators': [100, 200], 'max_depth': [None, 10,
↳20], 'min_samples_split': [2, 5]}

#base model
rf = RandomForestClassifier(random_state=42)

#grid search with 5-fold cross-validation
grid_search = GridSearchCV(
    estimator=rf,
    param_grid=param_grid,
```

```

cv=5,
scoring='roc_auc',
n_jobs=-1,
verbose=1)

#fit the grid search to the data
grid_search.fit(X, y)

```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```

[43]: GridSearchCV(cv=5, estimator=RandomForestClassifier(random_state=42), n_jobs=-1,
                param_grid={'max_depth': [None, 10, 20],
                            'min_samples_split': [2, 5],
                            'n_estimators': [100, 200]},
                scoring='roc_auc', verbose=1)

```

```

[45]: #print the best score and parameters
print("Best AUC Score:", grid_search.best_score_)
print("Best Parameters:", grid_search.best_params_)

```

Best AUC Score: 0.8479580156940711

Best Parameters: {'max_depth': 20, 'min_samples_split': 2, 'n_estimators': 200}

```

[47]: from sklearn.metrics import accuracy_score, roc_auc_score, classification_report
      from sklearn.model_selection import train_test_split

#split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y)
#fit on x train
grid_search.fit(X_train, y_train)

#get the best model from grid search
best_model = grid_search.best_estimator_

#predict on the test set
#for ROC AUC
y_pred = best_model.predict(X_test)
y_proba = best_model.predict_proba(X_test)[:, 1]

```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```

[49]: #print results of the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("AUC Score:", roc_auc_score(y_test, y_proba))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

```

Accuracy: 0.7538095238095238

AUC Score: 0.8458629988378008

Classification Report:

	precision	recall	f1-score	support
0	0.74	0.73	0.74	989
1	0.77	0.77	0.77	1111
accuracy			0.75	2100
macro avg	0.75	0.75	0.75	2100
weighted avg	0.75	0.75	0.75	2100

```
[ ]: #auc score of 85% means my model is learning from the data and the model is
      ↳ ranking winners over loses 85% of the time
      #we can see that blue(1) performanc is slightly higher than red(0) which is in
      ↳ line with expected results
      #this is mainly due a few reasons. First camera position giving a slight
      ↳ competitive advantage, first pick in draft
      #allows enemy team to be forced to ban the best champion in the current meta as
      ↳ blue can pick it up if its available
      #and better map and camera access to baron
      #seeing the model pick up the slight edge that blue has means the model is
      ↳ learning properly from the data
```

```
[ ]: #now trying XGBoost to compare model performance
      !pip install xgboost
```

```
[52]: from xgboost import XGBClassifier
      from sklearn.metrics import accuracy_score, roc_auc_score, classification_report
      from sklearn.model_selection import train_test_split

      #train/test split 80/20
      X_train, X_test, y_train, y_test = train_test_split(
          X, y, test_size=0.2, random_state=42, stratify=y)
      #initialize model
      xgb_model = XGBClassifier(
          n_estimators=300,
          max_depth=5,
          learning_rate=0.05,
          subsample=0.8,
          colsample_bytree=0.8,
          eval_metric='logloss',
          random_state=42)

      #train the model
      xgb_model.fit(X_train, y_train)
```



```
[52]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                  colsample_bylevel=None, colsample_bynode=None,
                  colsample_bytree=0.8, device=None, early_stopping_rounds=None,
                  enable_categorical=False, eval_metric='logloss',
                  feature_types=None, feature_weights=None, gamma=None,
                  grow_policy=None, importance_type=None,
                  interaction_constraints=None, learning_rate=0.05, max_bin=None,
                  max_cat_threshold=None, max_cat_to_onehot=None,
                  max_delta_step=None, max_depth=5, max_leaves=None,
                  min_child_weight=None, missing=nan, monotone_constraints=None,
                  multi_strategy=None, n_estimators=300, n_jobs=None,
                  num_parallel_tree=None, ...)
```

```
[55]: #predict
y_pred = xgb_model.predict(X_test)
y_proba = xgb_model.predict_proba(X_test)[:, 1]

#print model metrics
print("Accuracy:", accuracy_score(y_test, y_pred))
print("AUC Score:", roc_auc_score(y_test, y_proba))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.7690476190476191

AUC Score: 0.8570413158606054

Classification Report:

	precision	recall	f1-score	support
0	0.76	0.74	0.75	989
1	0.78	0.79	0.78	1111
accuracy			0.77	2100
macro avg	0.77	0.77	0.77	2100
weighted avg	0.77	0.77	0.77	2100

```
[ ]: #we can see a slightly higher performance on the model for the blue team
      #higher accuracy, AUC score, precision and recall
      #xgboost model is learning from mistakes in earlier trees
      #regularization from the XGBoost model is helping prevent overfitting
      #slightly higher metrics for blue(1) also shows model is learning and seeing
      ↳ blue win rate being
      #higher compared to red
```

```
[59]: #XGBoost model with GridSearchCV
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
```

```

param_grid = {
    'n_estimators': [200, 300],
    'max_depth': [4, 5, 6],
    'learning_rate': [0.03, 0.05, 0.1],
    'subsample': [0.8, 1.0],
    'colsample_bytree': [0.8, 1.0]}

#base model
xgb_base = XGBClassifier(
    eval_metric='logloss',
    random_state=42)

#grid search
xgb_grid = GridSearchCV(
    estimator=xgb_base,
    param_grid=param_grid,
    scoring='roc_auc',
    cv=5,
    verbose=1,
    n_jobs=-1)

#fit on training data only
xgb_grid.fit(X_train, y_train)

```

Fitting 5 folds for each of 72 candidates, totalling 360 fits

```

[59]: GridSearchCV(cv=5,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                          callbacks=None, colsample_bylevel=None,
                                          colsample_bynode=None,
                                          colsample_bytree=None, device=None,
                                          early_stopping_rounds=None,
                                          enable_categorical=False,
                                          eval_metric='logloss', feature_types=None,
                                          feature_weights=None, gamma=None,
                                          grow_policy=None, importance_type=None,
                                          interaction_constraint...,
                                          max_delta_step=None, max_depth=None,
                                          max_leaves=None, min_child_weight=None,
                                          missing=nan, monotone_constraints=None,
                                          multi_strategy=None, n_estimators=None,
                                          n_jobs=None, num_parallel_tree=None, ...),
                  n_jobs=-1,
                  param_grid={'colsample_bytree': [0.8, 1.0],
                              'learning_rate': [0.03, 0.05, 0.1],
                              'max_depth': [4, 5, 6], 'n_estimators': [200, 300],

```

```
        'subsample': [0.8, 1.0]},
    scoring='roc_auc', verbose=1)
```

```
[61]: #print best AUC score and best parameters
print("Best AUC Score:", xgb_grid.best_score_)
print("Best Parameters:", xgb_grid.best_params_)
```

Best AUC Score: 0.8529987720790245

Best Parameters: {'colsample_bytree': 0.8, 'learning_rate': 0.05, 'max_depth': 4, 'n_estimators': 200, 'subsample': 0.8}

```
[67]: #final model
XGBClassifier(
    n_estimators=200,
    max_depth=4,
    learning_rate=0.05,
    subsample=0.8,
    colsample_bytree=0.8,
    eval_metric='logloss',
    random_state=42)
```

```
[67]: XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=0.8, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric='logloss',
    feature_types=None, feature_weights=None, gamma=None,
    grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=0.05, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=4, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=200, n_jobs=None,
    num_parallel_tree=None, ...)
```

```
[77]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
#confusion matrix to show how well model performed, TN and TPs vs FN and FPs
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Red Win',
    ↪ 'Blue Win'])
disp.plot(cmap='Blues')
plt.title("Confusion Matrix")
plt.show()
```

